

Lab Topic 2 Block 1 Machine Learning

Hugo Knape & Zahra Jalil Pour & Niklas Larsson

11/12/2020

State of contribution

Assignment1: Zahra Jalilpour

Assignment2: Niklas Larsson

Assignment3: Hugo Knap

Assignment 1.

title: "Untitled" author: "z" date: '2020-11-20' output: html_document: default
pdf_document: default latex_engine: xelatex —

Assignment 1: LDA and logistic regression

In LDA, features in each class has multivariate normal distribution and common variance Σ . The covariance matrix is the same in all class $Cov(X) = \Sigma$. Random variable X is a vector $X = (X_1, X_2, \dots, X_p)$ Mean of each class is :

$$\hat{\mu}_k = \frac{1}{\#\{i ; y_i = k\}} \sum_{i ; y_i = k} x_i$$

Perior probabity or π_k is equal to:


$$\hat{\pi}_k = \frac{\#\{i ; y_i = k\}}{n}$$

X follows multivariate Gaussian distribution.

Task 1

Make a scatterplot

```
ggplot(iris, aes(x=Sepal.Width , y= Sepal.Length, shape= Species,  
color=Species)) + geom_point() +stat_ellipse()
```



Here we can see that we have more than two classes. Then LDA is the preferred linear classification technique, and LDA is a simple model for preparation and application. Here we can not use logistic regression, because we have more than two classes. Although we can use multi logistic regression but is rarely used for this purpose. LDA is a stable model when the classes are separated well. Also LDA generally is used for small data sets.

Task 2

2-a

Compute mean, covariance matrix and prior probabilities for each class:

```
## [1] "Covariance matrix in Setosa class"

##              Sepal.Length Sepal.Width
## Sepal.Length  0.12424898  0.09921633
## Sepal.Width   0.09921633  0.14368980

## [1] "Covariance matrix in Versicolor class"

##              Sepal.Length Sepal.Width
## Sepal.Length  0.26643265  0.08518367
## Sepal.Width   0.08518367  0.09846939

## [1] "Covariance matrix in Virginica class"
```

```
##           Sepal.Length Sepal.Width
## Sepal.Length  0.40434286  0.09376327
## Sepal.Width   0.09376327  0.10400408

## [1] "mean for each class"

##           setosa versicolor virginica
## Sepal.Length  5.006      5.936      6.588
## Sepal.Width   3.428      2.770      2.974

## [1] "prior probability for each class: "

## [1] 0.3333333
## [1] 0.3333333
## [1] 0.3333333
```

2-b

pooled covariance matrix:

```
## [1] "degree of freedom: "

##
##      setosa versicolor  virginica
##      49         49         49

## [1] "pooled covariance matrix: "

##           Sepal.Length Sepal.Width
## Sepal.Length  0.26500816  0.09272109
## Sepal.Width   0.09272109  0.11538776
```

2-c

probabilistic model: The probability of $p(X = x|Y = k)$ is given by:

$$f_k(x) = \frac{1}{(2\pi)^{p/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)\right)$$

By considering prior probability is $P(Y = k) = \pi_k$, and taking logarithm, we will find Linear discriminant function or Linear score function:

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log(\pi_k)$$

$$\hat{G}(x) = \arg \max_k \delta_k(x)$$

The decision boundary is the set of points in which two classes are equally probable:

$$\delta_k(x) = \delta_l(x)$$

We consider μ_k , σ_k by MLE as below:

$$\hat{\Sigma} = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T$$
$$\hat{\pi}_k = \frac{|\{i : y_i = k\}|}{n}$$

2_d

Compute discriminant functions for each class:

```
## [1] "Discriminant matrix for three different species"
```

```
##          setosa versicolor virginica
## [1,] 65.69175  58.09672  54.93699
## [2,] 53.22218  50.02352  46.49679
## [3,] 54.90090  47.79962  43.70231
## [4,] 51.69789  45.01644  40.69200
## [5,] 66.53111  56.98477  53.53975
## [6,] 77.32196  67.28187  64.77443
## [7,] 57.76144  47.52328  43.11160
## [8,] 62.48874  55.31354  51.92669
## [9,] 45.29187  39.45008  34.67140
## [10,] 55.24337  50.85914  47.30332
## [11,] 73.27960  65.61065  63.16137
## [12,] 60.12509  51.41841  47.51914
## [13,] 52.04036  48.07596  44.29302
## [14,] 46.13122  38.33812  33.27416
## [15,] 84.07045  75.90775  74.39605
## [16,] 90.97336  77.30264  75.41841
## [17,] 77.32196  67.28187  64.77443
## [18,] 65.69175  58.09672  54.93699
## [19,] 78.84626  72.28896  70.57921
## [20,] 71.75530  60.60356  57.35658
## [21,] 67.21605  63.10381  60.74177
## [22,] 69.73412  59.76795  56.55005
## [23,] 61.80380  49.19450  44.72466
## [24,] 61.64939  56.42550  53.32393
## [25,] 60.12509  51.41841  47.51914
## [26,] 54.40401  51.97109  48.70056
## [27,] 62.48874  55.31354  51.92669
## [28,] 66.87358  60.04429  57.14076
## [29,] 64.85240  59.20868  56.33423
## [30,] 54.90090  47.79962  43.70231
## [31,] 54.06154  48.91157  45.09955
## [32,] 67.21605  63.10381  60.74177
## [33,] 79.00068  65.05797  61.97995
## [34,] 84.56734  71.73628  69.39780
## [35,] 55.24337  50.85914  47.30332
```

##	[36,]	58.44638	53.64232	50.31362
##	[37,]	70.41906	65.88699	63.75208
##	[38,]	65.34928	55.03720	51.33598
##	[39,]	47.31305	40.28569	35.47793
##	[40,]	63.67057	57.26111	54.13046
##	[41,]	64.50993	56.14916	52.73322
##	[42,]	34.34660	36.38397	32.03598
##	[43,]	51.35542	41.95692	37.09099
##	[44,]	64.50993	56.14916	52.73322
##	[45,]	71.75530	60.60356	57.35658
##	[46,]	52.04036	48.07596	44.29302
##	[47,]	71.75530	60.60356	57.35658
##	[48,]	53.71907	45.85205	41.49854
##	[49,]	72.09777	63.66308	60.95760
##	[50,]	60.46756	54.47793	51.12015
##	[51,]	82.08291	92.59365	94.38906
##	[52,]	74.99195	80.90825	81.16643
##	[53,]	78.87990	89.81047	91.37875
##	[54,]	46.16487	55.85963	54.07370
##	[55,]	68.08905	79.51336	80.14407
##	[56,]	58.63443	63.93283	62.51390
##	[57,]	75.83131	79.79630	79.76919
##	[58,]	41.09509	45.00985	41.65760
##	[59,]	71.29206	82.29654	83.15437
##	[60,]	50.70412	53.35939	50.68851
##	[61,]	34.19218	43.61496	40.63524
##	[62,]	65.04045	69.49919	68.53450
##	[63,]	50.05282	64.76185	64.28602
##	[64,]	65.38292	72.55871	72.13552
##	[65,]	59.47379	62.82088	61.11666
##	[66,]	76.51625	85.91534	86.97121
##	[67,]	61.49497	63.65649	61.92319
##	[68,]	57.79508	65.04479	63.91114
##	[69,]	52.41647	68.65699	68.69357
##	[70,]	51.38906	59.47843	57.89053
##	[71,]	69.08282	71.17042	70.14757
##	[72,]	63.36174	71.72310	71.32898
##	[73,]	59.66185	73.11139	73.31693
##	[74,]	63.36174	71.72310	71.32898
##	[75,]	68.92840	78.40141	78.74683
##	[76,]	73.31324	83.13216	83.96091
##	[77,]	71.63453	85.35606	86.75539
##	[78,]	74.49507	85.07972	86.16468
##	[79,]	64.20110	70.61114	69.93174
##	[80,]	54.59207	62.26161	60.90083
##	[81,]	48.18605	56.69525	54.88023
##	[82,]	48.18605	56.69525	54.88023
##	[83,]	57.79508	65.04479	63.91114
##	[84,]	60.15873	68.93992	68.31868
##	[85,]	59.13132	59.76136	57.51565

##	[86,]	74.30701	74.78921	73.96440
##	[87,]	76.51625	85.91534	86.97121
##	[88,]	55.61948	71.44017	71.70387
##	[89,]	61.49497	63.65649	61.92319
##	[90,]	50.20723	57.53086	55.68676
##	[91,]	52.22841	58.36647	56.49329
##	[92,]	67.40411	73.39432	72.94205
##	[93,]	55.77389	64.20917	63.10461
##	[94,]	40.25573	46.12180	43.05484
##	[95,]	55.43142	61.14965	59.50359
##	[96,]	62.67680	65.60406	64.12696
##	[97,]	60.65562	64.76844	63.32043
##	[98,]	66.56475	74.50628	74.33929
##	[99,]	45.47992	49.74059	46.87167
##	[100,]	58.63443	63.93283	62.51390
##	[101,]	75.83131	79.79630	79.76919
##	[102,]	57.79508	65.04479	63.91114
##	[103,]	79.22237	92.86999	94.97977
##	[104,]	67.74658	76.45384	76.54306
##	[105,]	72.13141	81.18459	81.75713
##	[106,]	85.13151	102.60782	105.99862
##	[107,]	43.11627	45.84546	42.46413
##	[108,]	79.56484	95.92951	98.58078
##	[109,]	64.38915	80.90166	82.13202
##	[110,]	92.53130	99.83123	102.02273
##	[111,]	76.17378	82.85582	83.37020
##	[112,]	64.88604	76.73019	77.13377
##	[113,]	75.67689	87.02729	88.36845
##	[114,]	52.57089	61.42599	60.09430
##	[115,]	59.81626	65.88040	64.71767
##	[116,]	74.99195	80.90825	81.16643
##	[117,]	72.13141	81.18459	81.75713
##	[118,]	102.48280	111.24029	114.65465
##	[119,]	78.22860	101.21294	104.97627
##	[120,]	50.05282	64.76185	64.28602
##	[121,]	80.90109	90.64608	92.18529
##	[122,]	57.45261	61.98526	60.31013
##	[123,]	82.27097	102.88416	106.58933
##	[124,]	63.70421	74.78262	74.93000
##	[125,]	80.55861	87.58656	88.58427
##	[126,]	84.44657	96.48878	98.79660
##	[127,]	64.54357	73.67066	73.53276
##	[128,]	67.40411	73.39432	72.94205
##	[129,]	66.90722	77.56580	77.94030
##	[130,]	80.40420	94.81756	97.18354
##	[131,]	78.72549	97.04146	99.97802
##	[132,]	104.84645	115.13543	119.06219
##	[133,]	66.90722	77.56580	77.94030
##	[134,]	65.72539	75.61823	75.73653
##	[135,]	59.31938	70.05187	69.71592

```
## [136,] 86.31333 104.55539 108.20240
## [137,] 77.85249 80.63191 80.57572
## [138,] 72.97077 80.07264 80.35989
## [139,] 66.22228 71.44676 70.73828
## [140,] 78.87990 89.81047 91.37875
## [141,] 76.51625 85.91534 86.97121
## [142,] 78.87990 89.81047 91.37875
## [143,] 57.79508 65.04479 63.91114
## [144,] 79.71926 88.69852 89.98151
## [145,] 80.55861 87.58656 88.58427
## [146,] 74.49507 85.07972 86.16468
## [147,] 59.66185 73.11139 73.31693
## [148,] 72.13141 81.18459 81.75713
## [149,] 76.67066 78.68434 78.37195
## [150,] 65.04045 69.49919 68.53450
```

2-e

Compute equation of decision boundaries between classes:

```
## [1] decision boundary for Setosa ~ Versicolor
## [2]
## [3] W1 :

##           [,1]
## Sepal.Length -7.657399
## Sepal.Width 11.855698

## [1] "W0 : 5.15282159558004"

## [1] decision boundary for Versicolor ~ Verginica
## [2]
## [3] W1 :

##           [,1]
## Sepal.Length -2.5620509
## Sepal.Width 0.2908121

## [1] "W0 : 15.2083506778689"

## [1] decision boundary for Verginica ~ Setosa
## [2]
## [3] W1 :

##           [,1]
## Sepal.Length 10.21945
## Sepal.Width -12.14651


## [1] "W0 : -20.3611722734489"
```

3-a

predict function:

```
library(ggplot2)
D <- cbind(d1,d2,d3)

for(i in row(D)){
  predict <- apply(X = D[,], MARGIN = 1, FUN = which.max)
}
new_df <- iris
new_df <- cbind(new_df, predict)
new_df$predict[new_df$predict == "1"] <- "setosa"
new_df$predict[new_df$predict == "2"] <- "versicolor"
new_df$predict[new_df$predict == "3"] <- "virginica"
ggplot(new_df, aes(x=Sepal.Length , y= Sepal.Width, shape= predict,
color=predict)) + geom_point()+stat_ellipse()
```



```
Actual<- iris$Species
t_3 <- table(Actual, new_df$predict)
knitr::kable(t_3, caption = "MissClassification matrix for manual
prediction")
```

MissClassification matrix for manual prediction

	setosa	versicolor	virginica
--	--------	------------	-----------


```

setosa      49      1      0
versicolor  0      36     14
virginica   0      15     35
lda_error_3 = 1- sum(diag(t_3))/sum(t_3)
paste("MSE for manual LDA: ",lda_error_3)

## [1] "MSE for manual LDA: 0.2"

```

3_b

lda function

```

library(MASS)
new_data3c <- iris
fitted_lda <- lda(Species~Sepal.Length + Sepal.Width, data = new_data3c)
fitted_lda

## Call:
## lda(Species ~ Sepal.Length + Sepal.Width, data = new_data3c)
##
## Prior probabilities of groups:
##      setosa versicolor  virginica
## 0.3333333 0.3333333 0.3333333
##
## Group means:
##           Sepal.Length Sepal.Width
## setosa             5.006         3.428
## versicolor         5.936         2.770
## virginica          6.588         2.974
##
## Coefficients of linear discriminants:
##           LD1          LD2
## Sepal.Length -2.141178 -0.8152721
## Sepal.Width  2.768109 -2.0960764
##
## Proportion of trace:
##    LD1    LD2
## 0.9628 0.0372

#### Confusion Matrix####
Classification <- predict(fitted_lda, data = new_data3c)$class
Actual<- iris$Species
t <-table(Actual, Classification)
knitr::kable(t,caption = "MissClassification matrix for LDA")

```

MissClassification matrix for LDA

	setosa	versicolor	virginica
setosa	49	1	0

```

versicolor      0      36      14
virginica       0      15      35
lda_error = 1- sum(diag(t))/sum(t)
paste("MSE of LDA model: ", lda_error)

## [1] "MSE of LDA model: 0.2"

```


Miss classification matrix and Accuracy in both model are same.

```

library("mvtnorm")
bvn1 <- as.data.frame(mvrnorm(50, mu =m1, S ))
bvn2 <- as.data.frame(mvrnorm(50, mu=m2, S))
bvn3 <- as.data.frame(mvrnorm(50, mu=m3, S))
bv <- rbind(bvn1,bvn2,bvn3)

ggplot() +
  geom_point(data=bv, aes(bv$Sepal.Length, bv$Sepal.Width,
color=iris$Species))

```



By using new generate data, we can see the Setosa species is classified better than two classes, like other models. ### 5 Logistic Regression

```

library(nnet)
multi_model <- multinom(Species~Sepal.Length + Sepal.Width, data =iris)

## # weights: 12 (6 variable)
## initial value 164.791843

```

```

## iter 10 value 62.715967
## iter 20 value 59.808291
## iter 30 value 55.445984
## iter 40 value 55.375704
## iter 50 value 55.346472
## iter 60 value 55.301707
## iter 70 value 55.253532
## iter 80 value 55.243230
## iter 90 value 55.230241
## iter 100 value 55.212479
## final value 55.212479
## stopped after 100 iterations

multi_model

## Call:
## multinom(formula = Species ~ Sepal.Length + Sepal.Width, data = iris)
##
## Coefficients:
##             (Intercept) Sepal.Length Sepal.Width
## versicolor   -92.09924      40.40326   -40.58755
## virginica    -105.10096      42.30094   -40.18799
##
## Residual Deviance: 110.425
## AIC: 122.425

new_data <- iris
new_data$predicted<- predict(multi_model, new_data, type="class")
summary(new_data)

##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
##   Min.    :4.300   Min.    :2.000   Min.    :1.000   Min.    :0.100
##   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
##   Median :5.800   Median :3.000   Median :4.350   Median :1.300
##   Mean    :5.843   Mean    :3.057   Mean    :3.758   Mean    :1.199
##   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
##   Max.    :7.900   Max.    :4.400   Max.    :6.900   Max.    :2.500
##   Species      predicted
##   setosa      :50   setosa      :50
##   versicolor:50   versicolor:51
##   virginica  :50   virginica  :49
##
##
##

t_multi <- table(new_data$Species, new_data$predicted)
t_multi

##
##           setosa versicolor virginica
## setosa           50              0              0

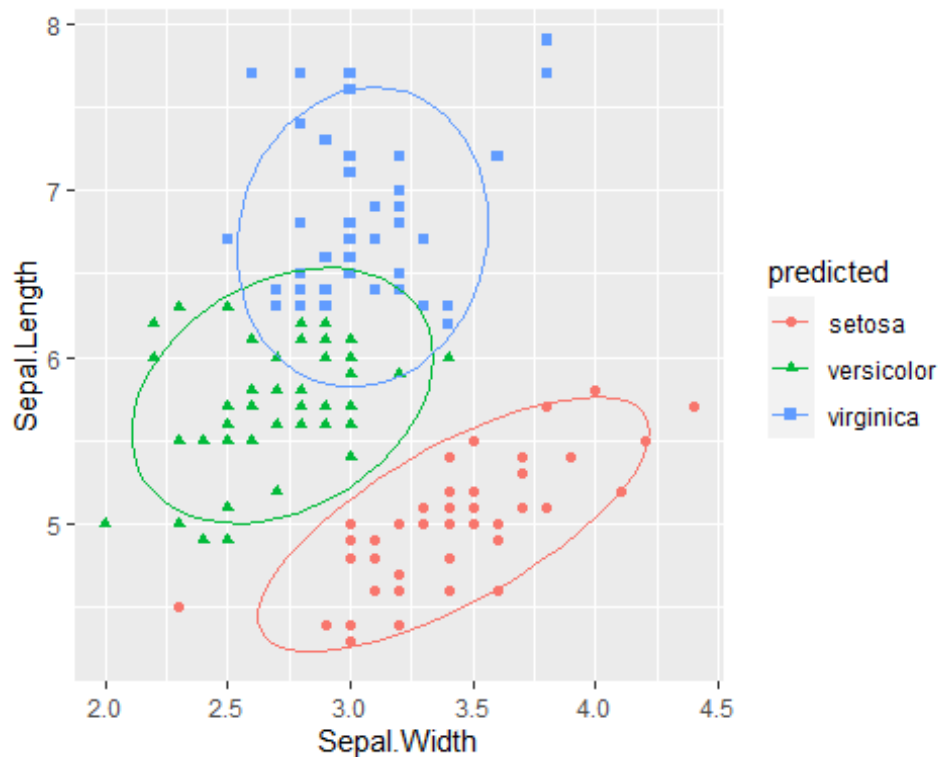
```

```
## versicolor      0      38      12
## virginica       0      13      37

ME_error <- 1-sum(diag(t_multi))/sum(t_multi)
paste("ME_error=", ME_error)

## [1] "ME_error= 0.166666666666667"

ggplot(new_data, aes(x=Sepal.Width , y= Sepal.Length, shape= predicted,
color=predicted)) + geom_point() +stat_ellipse()
```



We can see here that ME in logistic regression model is lower than two other models.

Assignment 2. Decision trees and Naïve Bayes for bank marketing

Task 1

Pre-process data by removing “Duration” feature, converting all character columns to categorical classes (factors) and split into train-, validation- and testset.

```
data = read.csv2("bank-full.csv")
data$duration = NULL

character_vars = lapply(data, class) == "character"
data[, character_vars] = lapply(data[, character_vars], as.factor)
#str(data)
```

```

n=dim(data)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.4))
train=data[id,]

id1=setdiff(1:n, id)
set.seed(12345)
id2=sample(id1, floor(n*0.3))
valid=data[id2,]

id3=setdiff(id1,id2)
test=data[id3,]

#sum(c(dim(train)[1],dim(valid)[1], dim(test)[1])) == dim(data)[1]

```

Task 2

Fit three different models from the given settings.

```

n = dim(train)[1]
fit_def = tree(formula = y ~., data=train, control = tree.control(nobs = n))
fit_node = tree(formula = y ~., data=train, control = tree.control(nobs = n,
minsize = 7000))
fit_dev = tree(formula = y ~., data=train, control = tree.control(nobs = n,
mindev = 0.0005))

misclass = matrix(0,3,2)
dimnames(misclass) = list( c("Default","Node","Deviance"),
c("Test","Validation") )

pred_def_train = predict(fit_def, newdata = train, type = "class")
pred_def_val = predict(fit_def, newdata = valid, type = "class")
misclass[1,1] = sum(train$y != pred_def_train)/length(train$y)
misclass[1,2] = sum(valid$y != pred_def_val)/length(valid$y)

pred_node_train = predict(fit_node, newdata = train, type = "class")
pred_node_val = predict(fit_node, newdata = valid, type = "class")
misclass[2,1] = sum(train$y != pred_node_train)/length(train$y)
misclass[2,2] = sum(valid$y != pred_node_val)/length(valid$y)

pred_dev_train = predict(fit_dev, newdata = train, type = "class")
pred_dev_val = predict(fit_dev, newdata = valid, type = "class")
misclass[3,1] = sum(train$y != pred_dev_train)/length(train$y)
misclass[3,2] = sum(valid$y != pred_dev_val)/length(valid$y)

```

```
misclass
```

```
##           Test Validation
## Default  0.10484406  0.1092679
## Node     0.10484406  0.1092679
## Deviance 0.09400575  0.1119221
```

From this result the default/nodesize trees performs best on validation data while deviance model does the opposite. Judging by these result the Deviance-model would be overfitting compared to the other two which also could be seen in the graphical illustrations below. The deviance model do have more potential to be optimized as this tree is much larger. The default and nodesize does only differ by one node which can explain why they perform identically!


Increasing the minimum node size results in a smaller tree as the data is divided in to larger areas (each area represent a terminal node) and there are less “classification squares” to decide a label between.

Decreasing the deviance forces the model to split nodes more often which will result in a larger tree as seen below. This does introduce more risk for overfitting. The deviance is calculated by


$$Deviance = - \sum_{i=1}^n p(c_i) \log(p(c_i))$$

where $p(c_i)$ is the probability of class c_i in the node. Using a lower value for deviance forces the probability for classification to be higher which results in more splits.


```
plot(fit_def, type = "uniform")
```



```
plot(fit_node, type = "uniform")
```



```
plot(fit_dev, type = "uniform")
```



Task 3


```
trainScore = rep(0,50)
testScore = rep(0,50)

for(i in 2:50){
  prunedTree = prune.tree(fit_dev, best = i)
  pred = predict(prunedTree, newdata = valid, type = "tree" )
  trainScore[i] = deviance(prunedTree)
  testScore[i] = deviance(pred)
}
```

As seen in the figure below the trees total deviance decreases rapidly during the first part as the model is able to distinguish between more features. As the number of leaves increases the training data's deviance keeps decreasing while the validation data starts to increase after around 20 leafs. This is where the optimal tree size is found. The reason why the training-deviance is larger than the validations is due to the training set being larger therefore more values to sum up.

```
plot(2:50, trainScore[2:50], type = "b", col = "red", ylim = c(8000,12000),
main = "Task 3", xlab = "Sequence", ylab = "Deviance")
points(2:50, testScore[2:50], type = "b", col = "blue")
legend("topright", c("Train","Validation"), pch = c("o","o"), col=
c("red","blue"))
```


Task 3



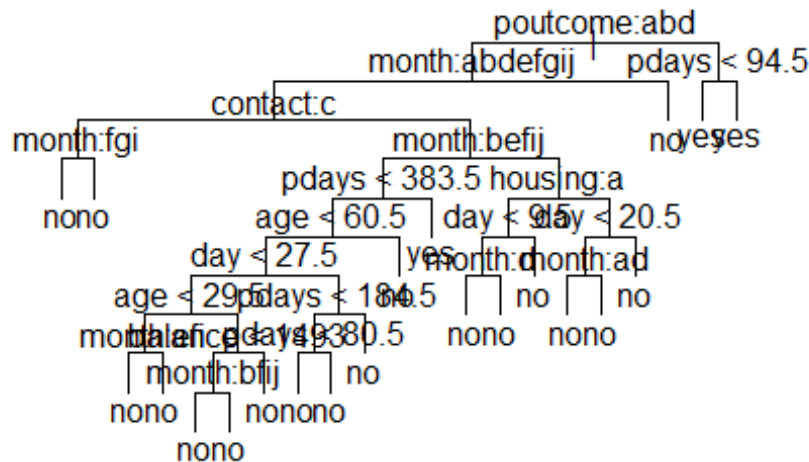
The optimal size (number of leaves/terminal nodes) is found where at the minimum deviance. The optimal number of leaves and the most significant features is shown in the summary of the optimal model:

```
opt_leaf = which.min(testScore[2:50])
opt_md1 = prune.tree(fit_dev, best = opt_leaf)
summary(opt_md1)

##
## Classification tree:
## snip.tree(tree = fit_dev, nodes = c(581L, 17L, 577L, 79L, 37L,
## 77L, 576L, 153L, 580L, 6L, 1157L, 16L, 5L, 1156L, 156L, 152L,
## 579L, 7L))
## Variables actually used in tree construction:
## [1] "outcome" "month" "contact" "pdays" "age" "day"
## "balance"
## [8] "housing"
## Number of terminal nodes: 21
## Residual mean deviance: 0.5706 = 10310 / 18060
## Misclassification error rate: 0.1041 = 1882 / 18084
```

The tree structure seems to use more terminal nodes than necessary as there are a lot of nodes which leads to the same label/ typ of terminal node, could be due to the deviance setting. The number of terminal nodes are indeed the same as the optimal which where found above. The outcome feature seems to be the most significant as it is set as the root node. This feature indicate wherever the customer was persuaded to subscribe a term deposit, which seems like a good starting point.

```
plot(opt_md1, type = "uniform")
text(opt_md1)
```



The prediction power of this model is ok, it does get the right answer in about 9/10 cases but its' main failure is by classifying False Negatives.

```
pred_valid = predict(opt_md1, newdata = valid, type = "class")
val_error = sum(pred_valid != valid$y)/length(valid$y)
val_conf_matrix = table(valid$y, pred_valid)
val_conf_matrix

##      pred_valid
##      no  yes
## no  11758  167
## yes  1319  319

sprintf("Misclassification rate: %.4f", val_error)

## [1] "Misclassification rate: 0.1096"
```

Task 4

By introducing a loss matrix to penalize the False Negative predictions the misclassification rate did increase but the number of False Negative classifications decreased! Using a loss matrix will direct the focus of learning towards the weighting of the matrix, in this case classifying False Negatives gives a five time larger penalty than False Positive and this will affect the overall classification which is why error rate increased.

Using the loss matrix does give a good result as the previous model have been trained on both the training and validation set while this model only have been exposed to the training set, suggesting that it learn faster (but more knowledge about the data is required to set up the loss matrix).

```
lossMatrix = matrix(c(0,5,1,0),2,2)

fit_matrix = rpart(formula = y~., data = train, parms = list(loss =
lossMatrix))
pred_matrix = predict(fit_matrix, newdata = test, type = "class")
error = sum(pred_matrix != test$y)/length(test$y)
cm_matrix = table(test$y, pred_matrix)
cm_matrix

##      pred_matrix
##      no  yes
## no 10880 1099
## yes  807  778

sprintf("Misclassification rate: %.4f", error)

## [1] "Misclassification rate: 0.1405"
```

Task 5

As can be seen in the graph below the optimal tree model seems to be the best model as the “area under curve” (AUC) is greater than the naive’s model. Comparing two models AUC could occasionally lead the wrong assumption within certain regions but in practice the AUC-measure performs well as a general comparison. The AUC is calculated using the trapezoidal rule. In this case the tree models ROC curve is always above the naive’s which also indicates that the tree model would perform best. Both models are well above the dashed line which could be interpret as a minimum boarder for a models predictive power, if a models is below this line the predictive power is worse than random guessing.

```
fit_naive = naiveBayes(formula = y~., data = train)

opt_pred = predict(opt_md1, newdata = test, type = "vector")
naive_pred = predict(fit_naive, newdata = test, type = "raw")

thres = seq(0, 0.95, 0.05)
TPR_tree = rep(0,length(thres))
FPR_tree = rep(0,length(thres))
TPR_naive = rep(0,length(thres))
FPR_naive = rep(0,length(thres))
y = test$y

for(i in 1:length(thres)){
  TP_tree = 0
  FP_tree = 0
```

```

TP_naive = 0
FP_naive = 0

for(j in 1:length(test$y)){

  if(opt_pred[j,2] >= thres[i]){
    if(y[j] == "yes"){
      TP_tree = TP_tree + 1
    }
    else{
      FP_tree = FP_tree + 1
    }
  }

  if(naive_pred[j,2] >= thres[i]){
    if(y[j] == "yes"){
      TP_naive = TP_naive + 1
    }
    else{
      FP_naive = FP_naive + 1
    }
  }

}

TPR_tree[i] = TP_tree/sum(test$y == "yes")
FPR_tree[i] = FP_tree/sum(test$y == "no")

TPR_naive[i] = TP_naive/sum(test$y == "yes")
FPR_naive[i] = FP_naive/sum(test$y == "no")
}

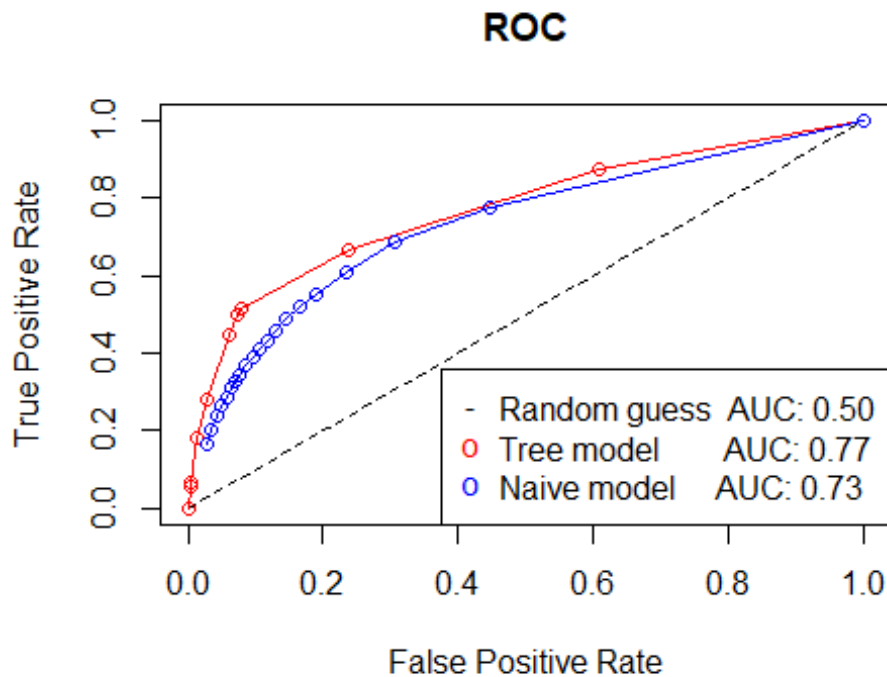
sub_model = seq(0,1,0.01)
AUC_tree = sum(abs(diff(FPR_tree)) * (head(TPR_tree,-1)+tail(TPR_tree,-1)))/2
AUC_naive = sum(abs(diff(FPR_naive)) * (head(TPR_naive,-1)+tail(TPR_naive,-1)))/2
AUC_sub = sum(abs(diff(sub_model)) * (head(sub_model,-1)+tail(sub_model,-1)))/2

plot(sub_model, sub_model, col = "black", type = "l", lty = 2, main = "ROC",
xlab = "False Positive Rate", ylab = "True Positive Rate")
lines(FPR_tree, TPR_tree, col = "red", type = "o")
lines(FPR_naive, TPR_naive, col = "blue", type = "o")

leg_text = c(sprintf("Random guess   AUC: %.2f", AUC_sub),sprintf("Tree model
AUC: %.2f", AUC_tree),sprintf("Naive model   AUC: %.2f", AUC_naive))

legend("bottomright", legend = leg_text, pch = c("-", "o", "o"), col =
c("black", "red", "blue"))

```



Assignment 3: Principal components for crime level analysis.

Task 1

```
communities <- read.csv("communities.csv")
communities <- communities[, -communities$ViolentCrimesPerPop]
#1
df <- (scale(communities))
s <- cov(df)
s.eigen <- eigen(s)
save <- c()
for (v in s.eigen$values) {
  h<- (v / sum(s.eigen$values))
  save <- c(save, h)
}
save[1]+save[2]

## [1] 0.4230255

nintyfive <- 0
i <- 1
while (nintyfive < 0.95) {
  nintyfive <- save[i] + nintyfive
  i <- i + 1
}
i-1
```

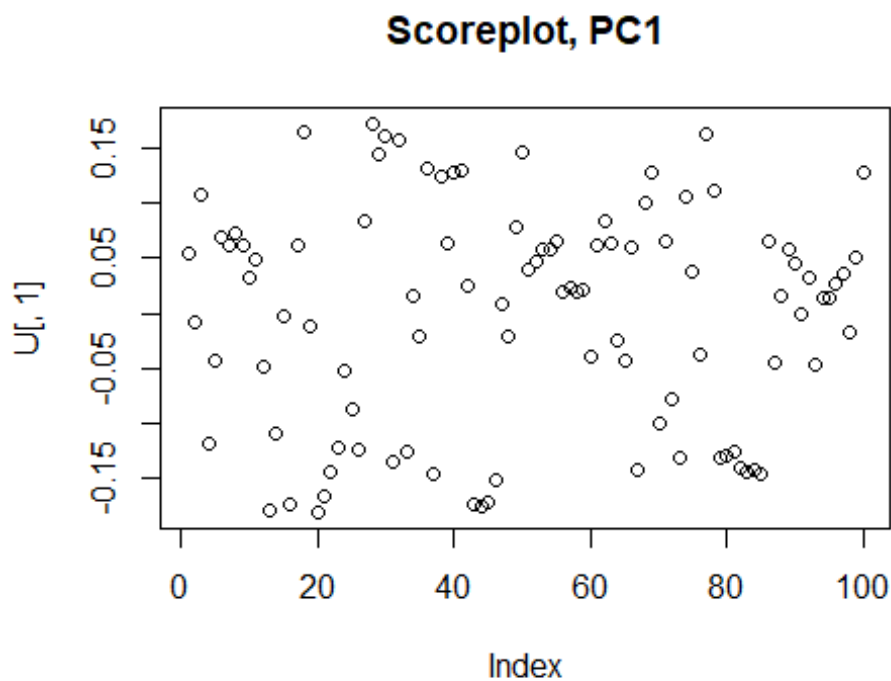
```
## [1] 35
```

The first two components explain around 42 percent of the variation. It needs 35 features to obtain at least 95% of variance in the data.

Task 2

```
#PRINCOMP
res <- princomp(df)
lambda <- res$sdev^2
#Plot
U <- res$loadings
x <- res$scores

plot(U[,1], main="Scoreplot, PC1")
```



The Scoreplot over PC1 shows that it seems around 15 features have a higher absolute value than 0.15, which means that they have a notable contribution to the first component.

The 5 features that have the biggest contribution to the first component are shown in the table down. We can see that many of these features have to do with Income and family, which is a very logical relationship to crime level.

```
#Absolute values
features <- as.data.frame(U[,1])
features$variable <- rownames(features)
features$`U[, 1]` <- abs(features$`U[, 1]`)
```


```
newdata <- features[order(features[,1],decreasing = TRUE ),]
newdata[1:5, ]
```

```
##           U[, 1]    variable
## medFamInc  0.1802364  medFamInc
## medIncome  0.1788053  medIncome
## PctKids2Par 0.1754391 PctKids2Par
## pctWInvInc 0.1736648  pctWInvInc
## PctFam2Par  0.1723885  PctFam2Par
```

The 5 features that has the biggest contribution to the first component are shown in the table above. We can see that many of they features has to do with Income and family which is a very logical relationship to crime level

#pc1 AGAINST PC2

```
pc12 <- as.data.frame(x[,1:2])
pc12$crimes <- communities$ViolentCrimesPerPop
library(ggplot2)
p1 <- ggplot(data = pc12, aes(x=Comp.1, y=Comp.2, color=crimes))+
  geom_point() +
  theme_bw() +
  labs(title = "PC1 vs PC2" , x = "PC1" , y = "PC2", colour =
"ViolentCrimesPerPop") +
  theme(axis.title.y = element_text(vjust = 0.5, size = 13 , face = "bold"))
+
  theme(axis.title.x = element_text(vjust = 0.5 ,size = 13 , face = "bold"))
+
  theme(plot.title = element_text(size = 14, face = "bold" , hjust = 0.4 )) +
  theme(axis.text.y = element_text(size = 11)) +
  theme(axis.text.x = element_text(size = 11))
p1
```




The plot above shows the PC1 versus PC2 with color of Violentcrimesperpop. The picture shows that the PC1 is pretty good to separate low and high values on violentcrimesperpop. It's hard to see any pattern between PC1 and PC2.

Task 3

```
model <- lm(crimes ~ poly(Comp.1,2), data = pc12)
pc12$pred <- (model$fitted.values)

ggplot(pc12, aes(x=Comp.1, y = crimes)) +
  geom_point(color= "orange") + theme_bw() +
  geom_line(aes(y=pred, lwd=1)) +
  labs(title = "PC1 vs Crimes" , x = "PC1" , y = "Crimes") +
  theme(axis.title.y = element_text(vjust = 0.5, size = 13 , face = "bold"))
+
  theme(axis.title.x = element_text(vjust = 0.5 ,size = 13 , face = "bold"))
+
  theme(plot.title = element_text(size = 14, face = "bold" , hjust = 0.4 )) +
  theme(axis.text.y = element_text(size = 11)) +
  theme(axis.text.x = element_text(size = 11))
```

The plot above shows the violentcrimesperpop versus PC1 with the fitted values from the model above as the black line. The graph shows that the model seems to capture the connection between the target and the feature pretty good.

Task 4

a)

```
library(boot)
data <- pc12
data2 <- data[order(data$Comp.1),]
mle <- lm(crimes ~ poly(Comp.1,2), data = data2)
rng <- function(data,mle2){
  n <- nrow(data)
  data1<- data.frame(crimes=data$crimes,Comp.1=data$Comp.1)
  data1$crimes <- rnorm(n,predict(mle2, newdata=data1),sd(mle2$residuals))
  return(data1)}

f1=function(data1){
  res=lm(crimes~poly(Comp.1,degree=2),data1)
  ViolentCrimesPerPopP=predict(res,newdata=data2)
  return(ViolentCrimesPerPopP)}


fit <- lm(crimes ~ poly(Comp.1,2), data = data2)
crimesP=predict(fit)
```

```

res2 <- boot(data2, statistic=f1, R=3000, mle = fit, ran.gen=rng ,
sim="parametric")
e2 <- envelope(res2)

plot(data$Comp.1,data$crimes , pch=21, bg="orange", ylim=c(-0.3,1.1), main =
"Confidence interval")
points(data2$Comp.1,crimesP,type="l") #plot fitted line
#plot confidence bands
points(data2$Comp.1,e2$point[2,], type="l", col="blue")
points(data2$Comp.1,e2$point[1,], type="l", col="blue")

```



The plot above shows the plot from task 3 but now with confidence intervals. You can see that the confidence intervals are wider in the beginning and the end. That is because there are fewer points there compared to in the middle of the plot which gives them more uncertainty.

b)

```

f1 <- function(data1){
  res <- lm(crimes ~ poly(Comp.1,2), data = data1)
  crimesP <- predict(res, newdata = data2 )
  n <- length(data2$crimes)
  predictedP <- rnorm(n, crimesP, sd(mle$residuals))
  return(predictedP)
}

```

```

rng <- function(data, mle ){
  data1 <- data.frame(crimes=data$crimes, Comp.1=data$Comp.1)
  n <- length(data$crimes)
  data1$Crimes <- rnorm(n, predict(mle, newdata = data1), sd(mle$residuals))
  return(data1)
}

```

```

res <- boot(data2, statistic=f1, R=3000,
            mle=mle,ran.gen=rng, sim="parametric")

```

```

e <- envelope(res) #compute prediction bands


```

Warning in envelope(res): unable to achieve requested overall error rate

```

fit <- lm(crimes ~ poly(Comp.1,2), data = data2)
crimesP=predict(fit)
plot(data$Comp.1,data$crimes , pch=21, bg="orange", ylim=c(-0.3,1.1))
points(data2$Comp.1,crimesP,type="l") #plot fitted line
#plot prediction bands
points(data2$Comp.1,e$point[2,], type="l", col="blue")
points(data2$Comp.1,e$point[1,], type="l", col="blue")

```



The plot above shows the plot from task 3 but now with prediction intervals. You can see that the predictions bands have almost the same width all over the plot and that the intervals alot of the points in the plot. This is because prediction is made out of every fitted point instead of confidence interval which is made out of the fitted line.

Appendix:

```
knitr::opts_chunk$set(echo = TRUE)
library(kknn)
library(ggplot2)
library(e1071)
library(tree)
library(rpart)
library(ggplot2)
library(datasets)

ggplot(iris, aes(x=Sepal.Width , y= Sepal.Length, shape= Species,
color=Species)) + geom_point() +stat_ellipse()

data(iris)
### covariance matrix per each class
paste("Covariance matrix in Setosa class")
(var(subset(iris,subset=Species=='setosa',select=c(1:2))) -> S1)
paste("Covariance matrix in Versicolor class")
(var(subset(iris,subset=Species=='versicolor',select=c(1:2))) -> S2)
paste("Covariance matrix in Virginica class")
(var(subset(iris,subset=Species=='virginica',select=c(1:2))) -> S3)

### mean per each class

m1 <- colMeans(subset(iris,subset=Species=='setosa',select=c(1:2)))
m2 <- colMeans(subset(iris,subset=Species=='versicolor',select=c(1:2)))
m3 <- colMeans(subset(iris,subset=Species=='virginica',select=c(1:2)))

paste("mean for each class")
results <- cbind(m1,m2,m3)
colnames(results) <- c("setosa","versicolor","virginica")
rownames(results) <- c("Sepal.Length", "Sepal.Width")
results

### prior probability per each class
paste("prior probability for each class: ")
(nrow(subset(iris,subset=Species=='setosa')) / nrow(iris) -> prio_prob1)
(nrow(subset(iris,subset=Species=='versicolor')) / nrow(iris) -> prio_prob2)
```

```

(nrow(subset(iris,subset=Species=='virginica')) / nrow(iris) -> prio_prob3)

paste("degree of freedom: ")
(table(iris$Species)-1 -> dof)

paste("pooled covariance matrix: ")
((dof[1]*S1+dof[2]*S2+dof[3]*S3)/sum(dof) -> S)
multi_gusi_model <- function(x,mean,sigma){
  p <- ncol(x)
  sig_inv <- solve(sigma)
  state1 <- (sqrt((2*pi)^p) *(det(sigma)))
  state2 <- -(1/2)*(as.matrix(x-mean) %*% solve(sigma) %*% t(x-mean) )
  model <- (1/state1) * exp(state2)
  return(model)
}

discriminant <- function(x,mean_i, pooled_cov, prior_i){
  #mu <- as.matrix(mean_i)
  #browser()
  state1 <- as.matrix(x) %*% solve(pooled_cov) %*% mean_i
  state2 <- 0.5 * t(mean_i) %*% solve(pooled_cov) %*% mean_i
  d <- state1 - as.numeric(state2) + as.numeric(log(prior_i))
  return(d)
}
#example for one class
#pooled <- ((dof[1]*S1+dof[2]*S2+dof[3]*S3)/sum(dof) -> S)
prior <- (nrow(subset(iris,subset=Species=='setosa')) / nrow(iris) ->
prio_prob1)
d1 <-discriminant(iris[c(1,2)] ,matrix(m1,ncol = 1),S,prio_prob1)

d2 <-discriminant(iris[c(1,2)] ,matrix(m2,ncol = 1),S,prio_prob2)

d3 <-discriminant(iris[c(1,2)] ,matrix(m3,ncol = 1),S,prio_prob3)

D_mat <- cbind(d1,d2,d3)
colnames(D_mat) <- c("setosa","versicolor","virginica")
paste("Discriminant matrix for three different species" )
D_mat
w0 <- function(pi_k, pi_l, mu_k, mu_l, pooled_cov){
  m1 <- as.matrix(m1)
  m2 <- as.matrix(m2)
  w0 <- -(1/2)*t(mu_k + mu_l) %*% solve(pooled_cov) %*% (mu_k - mu_l)
  return(w0)
}
w1 <-function(pi_k, pi_l, mu_k, mu_l, pooled_cov){
  m1 <- as.matrix(m1)
  m2 <- as.matrix(m2)
  w1 <- solve(pooled_cov) %*% (mu_k - mu_l)

```

```

    return(as.matrix(w1))
}
msg1 <- paste("decision boundary for Setosa ~ Versicolor\n\n")
msg2 <- paste(msg1, "W1 : ")
noquote(strsplit(msg2, "\n")[[1]])
w1(0.33,0.33,m1,m2,S)
paste("W0 : ", w0(0.33,0.33,m1,m2,S))

msg1 <- paste("decision boundary for Versicolor ~ Virginica\n\n")
msg2 <- paste(msg1, "W1 : ")
noquote(strsplit(msg2, "\n")[[1]])
w1(0.33,0.33,m2,m3,S)
paste("W0 : ",w0(0.33,0.33,m2,m3,S))

msg1 <- paste("decision boundary for Virginica ~ Setosa\n\n")
msg2 <- paste(msg1, "W1 : ")
noquote(strsplit(msg2, "\n")[[1]])
w1(0.33,0.33,m3,m1,S)
paste("W0 : ", w0(0.33,0.33,m3,m1,S))
library(ggplot2)
D <- cbind(d1,d2,d3)

for(i in row(D)){
  predict <- apply(X = D[,], MARGIN = 1, FUN = which.max)
}
new_df <- iris
new_df <- cbind(new_df, predict)
new_df$predict[new_df$predict == "1"] <- "setosa"
new_df$predict[new_df$predict == "2"] <- "versicolor"
new_df$predict[new_df$predict == "3"] <- "virginica"
ggplot(new_df, aes(x=Sepal.Length , y= Sepal.Width, shape= predict,
color=predict)) + geom_point()+stat_ellipse()

Actual<- iris$Species
t_3 <-table(Actual, new_df$predict)
knitr::kable(t_3, caption = "MissClassification matrix for manual
prediction")
lda_error_3 = 1- sum(diag(t_3))/sum(t_3)
paste("MSE for manual LDA: " ,lda_error_3)
library(MASS)
new_data3c <- iris
fitted_lda <- lda(Species~Sepal.Length + Sepal.Width, data = new_data3c)
fitted_lda
#### Confusion Matrix####
Classification <- predict(fitted_lda, data = new_data3c)$class
Actual<- iris$Species
t <-table(Actual, Classification)
knitr::kable(t,caption = "MissClassification matrix for LDA")
lda_error = 1- sum(diag(t))/sum(t)

```

```

paste("MSE of LDA model: ", lda_error)

library("mvtnorm")
bvn1 <- as.data.frame(mvrnorm(50, mu =m1, S ))
bvn2 <- as.data.frame(mvrnorm(50, mu=m2, S))
bvn3 <- as.data.frame(mvrnorm(50, mu=m3, S))
bv <- rbind(bvn1,bvn2,bvn3)

ggplot() +
  geom_point(data=bv, aes(bv$Sepal.Length, bv$Sepal.Width,
color=iris$Species))
library(nnet)
multi_model <- multinom(Species~Sepal.Length + Sepal.Width, data =iris)
multi_model
new_data <- iris
new_data$predicted<- predict(multi_model, new_data, type="class")
summary(new_data)
t_multi <- table(new_data$Species, new_data$predicted)
t_multi
ME_error <- 1-sum(diag(t_multi))/sum(t_multi)
paste( "ME_error= ", ME_error)

ggplot(new_data, aes(x=Sepal.Width , y= Sepal.Length, shape= predicted,
color=predicted)) + geom_point() +stat_ellipse()

data = read.csv2("bank-full.csv")
data$duration = NULL

character_vars = lapply(data, class) == "character"
data[, character_vars] = lapply(data[, character_vars], as.factor)
#str(data)

n=dim(data)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.4))
train=data[id,]

id1=setdiff(1:n, id)
set.seed(12345)
id2=sample(id1, floor(n*0.3))
valid=data[id2,]

id3=setdiff(id1,id2)
test=data[id3,]

#sum(c(dim(train)[1],dim(valid)[1], dim(test)[1])) == dim(data)[1]

```

```

n = dim(train)[1]
fit_def = tree(formula = y ~., data=train, control = tree.control(nobs = n))
fit_node = tree(formula = y ~., data=train, control = tree.control(nobs = n,
minsize = 7000))
fit_dev = tree(formula = y ~., data=train, control = tree.control(nobs = n,
mindev = 0.0005))

misclass = matrix(0,3,2)
dimnames(misclass) = list( c("Default","Node","Deviance"),
c("Test","Validation") )

pred_def_train = predict(fit_def, newdata = train, type = "class")
pred_def_val = predict(fit_def, newdata = valid, type = "class")
misclass[1,1] = sum(train$y != pred_def_train)/length(train$y)
misclass[1,2] = sum(valid$y != pred_def_val)/length(valid$y)

pred_node_train = predict(fit_node, newdata = train, type = "class")
pred_node_val = predict(fit_node, newdata = valid, type = "class")
misclass[2,1] = sum(train$y != pred_node_train)/length(train$y)
misclass[2,2] = sum(valid$y != pred_node_val)/length(valid$y)

pred_dev_train = predict(fit_dev, newdata = train, type = "class")
pred_dev_val = predict(fit_dev, newdata = valid, type = "class")
misclass[3,1] = sum(train$y != pred_dev_train)/length(train$y)
misclass[3,2] = sum(valid$y != pred_dev_val)/length(valid$y)

misclass

plot(fit_def, type = "uniform")
plot(fit_node, type = "uniform")
plot(fit_dev, type = "uniform")

trainScore = rep(0,50)
testScore = rep(0,50)

for(i in 2:50){
  prunedTree = prune.tree(fit_dev, best = i)
  pred = predict(prunedTree, newdata = valid, type = "tree" )
  trainScore[i] = deviance(prunedTree)
  testScore[i] = deviance(pred)
}

plot(2:50, trainScore[2:50], type = "b", col = "red", ylim = c(8000,12000),
main = "Task 3", xlab = "Sequence", ylab = "Deviance")
points(2:50, testScore[2:50], type = "b", col = "blue")

```



```
legend("topright", c("Train","Validation"), pch = c("o","o"), col=
c("red","blue"))
```

```
opt_leaf = which.min(testScore[2:50])
opt_mdl = prune.tree(fit_dev, best = opt_leaf)
summary(opt_mdl)
plot(opt_mdl, type = "uniform")
text(opt_mdl)
pred_valid = predict(opt_mdl, newdata = valid, type = "class")
val_error = sum(pred_valid != valid$y)/length(valid$y)
val_conf_matrix = table(valid$y, pred_valid)
val_conf_matrix
sprintf("Misclassification rate: %.4f", val_error)
```

```
lossMatrix = matrix(c(0,5,1,0),2,2)
```

```
fit_matrix = rpart(formula = y~., data = train, parms = list(loss =
lossMatrix))
pred_matrix = predict(fit_matrix, newdata = test, type = "class")
error = sum(pred_matrix != test$y)/length(test$y)
cm_matrix = table(test$y, pred_matrix)
cm_matrix
sprintf("Misclassification rate: %.4f", error)
```

```
fit_naive = naiveBayes(formula = y~., data = train)
```

```
opt_pred = predict(opt_mdl, newdata = test, type = "vector")
naive_pred = predict(fit_naive, newdata = test, type = "raw")
```

```
thres = seq(0, 0.95, 0.05)
TPR_tree = rep(0,length(thres))
FPR_tree = rep(0,length(thres))
TPR_naive = rep(0,length(thres))
FPR_naive = rep(0,length(thres))
y = test$y
```

```
for(i in 1:length(thres)){
  TP_tree = 0
  FP_tree = 0
  TP_naive = 0
  FP_naive = 0

  for(j in 1:length(test$y)){
    if(opt_pred[j,2] >= thres[i]){
      if(y[j] == "yes"){
```

```

        TP_tree = TP_tree + 1
      }
    else{
      FP_tree = FP_tree + 1
    }
  }

  if(naive_pred[j,2] >= thres[i]){
    if(y[j] == "yes"){
      TP_naive = TP_naive + 1
    }
    else{
      FP_naive = FP_naive + 1
    }
  }
}

TPR_tree[i] = TP_tree/sum(test$y == "yes")
FPR_tree[i] = FP_tree/sum(test$y == "no")

TPR_naive[i] = TP_naive/sum(test$y == "yes")
FPR_naive[i] = FP_naive/sum(test$y == "no")
}

sub_model = seq(0,1,0.01)
AUC_tree = sum(abs(diff(FPR_tree)) * (head(TPR_tree,-1)+tail(TPR_tree,-1)))/2
AUC_naive = sum(abs(diff(FPR_naive)) * (head(TPR_naive,-1)+tail(TPR_naive,-1)))/2
AUC_sub = sum(abs(diff(sub_model)) * (head(sub_model,-1)+tail(sub_model,-1)))/2

plot(sub_model, sub_model, col = "black", type = "l", lty = 2, main = "ROC",
xlab = "False Positive Rate", ylab = "True Positive Rate")
lines(FPR_tree, TPR_tree, col = "red", type = "o")
lines(FPR_naive, TPR_naive, col = "blue", type = "o")

leg_text = c(sprintf("Random guess   AUC: %.2f", AUC_sub),sprintf("Tree model
AUC: %.2f", AUC_tree),sprintf("Naive model   AUC: %.2f", AUC_naive))

legend("bottomright", legend = leg_text, pch = c("-", "o", "o"), col =
c("black", "red", "blue"))

communities <- read.csv("communities.csv")
communities <- communities[, -communities$ViolentCrimesPerPop]
#1
df <- (scale(communities))
s <- cov(df)
s.eigen <- eigen(s)

```

```

save <- c()
for (v in s.eigen$values) {
  h<- (v / sum(s.eigen$values))
  save <- c(save, h)
}
save[1]+save[2]

nintyfive <- 0
i <- 1
while (nintyfive < 0.95) {
  nintyfive <- save[i] + nintyfive
  i <- i + 1
}
i-1
#PRINCOMP
res <- princomp(df)
lambda <- res$sdev^2
#Plot
U <- res$loadings
x <- res$scores

plot(U[,1], main="Scoreplot, PC1")
#Absolute values
features <- as.data.frame(U[,1])
features$variable <- rownames(features)
features$`U[, 1]` <- abs(features$`U[, 1]`)
newdata <- features[order(features[,1],decreasing = TRUE ),]
newdata[1:5, ]
#pc1 AGAINST PC2

pc12 <- as.data.frame(x[,1:2])
pc12$crimes <- communities$ViolentCrimesPerPop
library(ggplot2)
p1 <- ggplot(data = pc12, aes(x=Comp.1, y=Comp.2, color=crimes))+
  geom_point() +
  theme_bw() +
  labs(title = "PC1 vs PC2" , x = "PC1" , y = "PC2", colour =
"ViolentCrimesPerPop") +
  theme(axis.title.y = element_text(vjust = 0.5, size = 13 , face = "bold"))
+
  theme(axis.title.x = element_text(vjust = 0.5 ,size = 13 , face = "bold"))
+
  theme(plot.title = element_text(size = 14, face = "bold" , hjust = 0.4 )) +
  theme(axis.text.y = element_text(size = 11)) +
  theme(axis.text.x = element_text(size = 11))
p1
model <- lm(crimes ~ poly(Comp.1,2), data = pc12)
pc12$pred <- (model$fitted.values)

```

```

ggplot(pc12, aes(x=Comp.1, y = crimes)) +
  geom_point(color= "orange") + theme_bw() +
  geom_line(aes(y=pred, lwd=1)) +
  labs(title = "PC1 vs Crimes" , x = "PC1" , y = "Crimes") +
  theme(axis.title.y = element_text(vjust = 0.5, size = 13 , face = "bold"))
+
  theme(axis.title.x = element_text(vjust = 0.5 ,size = 13 , face = "bold"))
+
  theme(plot.title = element_text(size = 14, face = "bold" , hjust = 0.4 )) +
  theme(axis.text.y = element_text(size = 11)) +
  theme(axis.text.x = element_text(size = 11))

library(boot)
data <- pc12
data2 <- data[order(data$Comp.1),]
mle <- lm(crimes ~ poly(Comp.1,2), data = data2)
rng <- function(data,mle2){
  n <- nrow(data)
  data1<- data.frame(crimes=data$crimes,Comp.1=data$Comp.1)
  data1$crimes <- rnorm(n,predict(mle2, newdata=data1),sd(mle2$residuals))
  return(data1)}

f1=function(data1){
  res=lm(crimes~poly(Comp.1,degree=2),data1)
  ViolentCrimesPerPopP=predict(res,newdata=data2)
  return(ViolentCrimesPerPopP)}

fit <- lm(crimes ~ poly(Comp.1,2), data = data2)
crimesP=predict(fit)

res2 <- boot(data2, statistic=f1, R=3000, mle = fit, ran.gen=rng ,
sim="parametric")
e2 <- envelope(res2)

plot(data$Comp.1,data$crimes , pch=21, bg="orange", ylim=c(-0.3,1.1), main =
"Confidence interval")
points(data2$Comp.1,crimesP,type="l") #plot fitted line
#plot confidence bands
points(data2$Comp.1,e2$point[2,], type="l", col="blue")
points(data2$Comp.1,e2$point[1,], type="l", col="blue")

f1 <- function(data1){
  res <- lm(crimes ~ poly(Comp.1,2), data = data1)
  crimesP <- predict(res, newdata = data2 )
  n <- length(data2$crimes)
  predictedP <- rnorm(n, crimesP, sd(mle$residuals))
  return(predictedP)
}

```

```

}

rng <- function(data, mle ){
  data1 <- data.frame(crimes=data$crimes, Comp.1=data$Comp.1)
  n <- length(data$crimes)
  data1$Crimes <- rnorm(n, predict(mle, newdata = data1), sd(mle$residuals))
  return(data1)
}

res <- boot(data2, statistic=f1, R=3000,
            mle=mle,ran.gen=rng, sim="parametric")

e <- envelope(res) #compute prediction bands
fit <- lm(crimes ~ poly(Comp.1,2), data = data2)
crimesP=predict(fit)
plot(data$Comp.1,data$crimes , pch=21, bg="orange", ylim=c(-0.3,1.1))
points(data2$Comp.1,crimesP,type="l") #plot fitted line
#plot prediction bands
points(data2$Comp.1,e$point[2,], type="l", col="blue")
points(data2$Comp.1,e$point[1,], type="l", col="blue")

```